

# Machine Learning Essentials SS25 - Exercise Sheet 6

## Instructions

- `TODO` 's indicate where you need to complete the implementations.
- You may use external resources, but **write your own solutions**.
- Provide concise, but comprehensible comments to explain what your code does.
- Code that's unnecessarily extensive and/or not well commented will not be scored.

```
In [1]: import matplotlib.pyplot as plt
import numpy as np
import torch as tc
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset
from torch.utils.data import DataLoader
from torchsummary import summary

np.random.seed(42)
tc.manual_seed(42)

device = tc.device("cuda" if tc.cuda.is_available() else "cpu")
```

## Exercise 2 - CNN Classifier

The SIGNS dataset is a collection of 6 signs representing numbers from 0 to 5. We first load the data and have the shapes printed out. The split into train, validation and test set has already been carried out.

```
In [2]: # Load the dataset
X_train = np.load('sign_data/X_train.npy')
Y_train = np.load('sign_data/Y_train.npy')
X_val = np.load('sign_data/X_val.npy')
Y_val = np.load('sign_data/Y_val.npy')
X_test = np.load('sign_data/X_test.npy')
Y_test = np.load('sign_data/Y_test.npy')

# print the shape of the dataset
print("X_train shape: " + str(X_train.shape))
print("Y_train shape: " + str(Y_train.shape))
print("X_val shape: " + str(X_val.shape))
print("Y_val shape: " + str(Y_val.shape))
print("X_test shape: " + str(X_test.shape))
print("Y_test shape: " + str(Y_test.shape)+"\n")
print("classes: " + str(np.unique(Y_train)))

# check if classes are balanced
print("Counts of classes in Y_train: " + str(np.unique(Y_train, return_counts=True)))
print("Counts of classes in Y_val: " + str(np.unique(Y_val, return_counts=True)))
print("Counts of classes in Y_test: " + str(np.unique(Y_test, return_counts=True)))
```

```

X_train shape: (960, 64, 64, 3)
Y_train shape: (960,)
X_val shape: (120, 64, 64, 3)
Y_val shape: (120,)
X_test shape: (120, 64, 64, 3)
Y_test shape: (120,)

```

```

classes: [0 1 2 3 4 5]
Counts of classes in Y_train: [160 160 160 160 160 160]
Counts of classes in Y_val: [20 20 20 20 20 20]
Counts of classes in Y_test: [20 20 20 20 20 20]


```

The classes are balanced so that accuracy is an appropriate measure for evaluating a classifier. We next visualize an instance of each class.

```

In [3]: fig, axes = plt.subplots(1, 6, figsize=(20, 10))
        for i in range(6):
            # get indices where the label is i
            idx = np.where(Y_train == i)[0][0]
            axes[i].imshow(X_train[idx])
            axes[i].set_title("y = " + str(i))

```



Pixels in each channel (RGB) of the images take values in the range [0, 255]. However, it is desirable to have absolute values in the range [0, 1] as input for neural network architectures to avoid exploding or vanishing gradient problems. Through the following cell, we apply a simple data scaling procedure: we divide the values of the pixels by 255. As an alternative, you can use the `StandardScaler()` function of the scikit-learn library.

```

In [4]: X_train = X_train/255
        X_val = X_val/255
        X_test = X_test/255

```

## Task 1

Use pytorch to build the model. Take a look at the [documentation](#) for an introduction, a detailed tutorial, for example for classifiers, can be found [here](#).

Implement the following architecture:

- Conv2d: 4 output channels, 3 by 3 filter size, stride 1, padding "same"
- BatchNorm2d: 4 output channels
- ReLU activation
- MaxPool2d: 2 by 2 filter size, stride 2, padding 0
- Conv2d: 8 output channels, 3 by 3 filter size, stride 1, padding "same"
- BatchNorm2d: 8 output channels
- ReLU activation

- MaxPool2d: Use a 2 by 2 filter size, stride 2, padding 0
- Flatten the previous output
- Linear: 64 output neurons
- ReLu activation function
- Linear: 6 output neurons
- LogSoftmax

We use the [LogSoftmax](#) here instead of the Softmax for computational reasons.

Accordingly, the loss function is not CrossEntropyLoss but NLLLoss. When flattening, be careful not to do it with the batch dimension but only with the height, width and channel dimension.

```
In [5]: class CNN_Classifier(nn.Module):
    def __init__(self):
        super().__init__()

        # TODO: Initialize the layers of the CNN

        # First convolutional block
        self.conv1 = nn.Conv2d(3, 4, kernel_size=3, stride=1, padding='same')
        self.bn1 = nn.BatchNorm2d(4)
        self.relu1 = nn.ReLU()
        self.pool1 = nn.MaxPool2d(kernel_size=2, stride=2, padding=0)

        # Second convolutional block
        self.conv2 = nn.Conv2d(4, 8, kernel_size=3, stride=1, padding='same')
        self.bn2 = nn.BatchNorm2d(8)
        self.relu2 = nn.ReLU()
        self.pool2 = nn.MaxPool2d(kernel_size=2, stride=2, padding=0)

        # Flatten layer
        self.flatten = nn.Flatten()

        # Fully connected layers
        # After two 2x2 max pooling operations, 64x64 -> 32x32 -> 16x16
        # So we have 8 channels * 16 * 16 = 2048 features
        self.fc1 = nn.Linear(8 * 16 * 16, 64)
        self.relu3 = nn.ReLU()
        self.fc2 = nn.Linear(64, 6)
        self.logsoftmax = nn.LogSoftmax(dim=1)

    def forward(self, X):
        # TODO: Implement the forward pass
        # First convolutional block
        X = self.conv1(X)
        X = self.bn1(X)
        X = self.relu1(X)
        X = self.pool1(X)

        # Second convolutional block
        X = self.conv2(X)
        X = self.bn2(X)
        X = self.relu2(X)
        X = self.pool2(X)

        # Flatten and fully connected layers
        X = self.flatten(X)
```

```

X = self.fc1(X)
X = self.relu3(X)
X = self.fc2(X)
X = self.logsoftmax(X)

return X

```

To test your model you can forward some random numbers. The shape of the output should be (2, 6).

```

In [6]: cnn_model = CNN_Classifier()
# dummy sample of batch size 2
X_random = tc.randn(2, 3, 64, 64)
output = cnn_model(X_random)

print("Output shape: " + str(output.shape))

```

Output shape: torch.Size([2, 6])

torchsummary.summary provides a nice overview of the model and the number of learnable parameters:

```

In [7]: summary(cnn_model, input_size=(3, 64, 64), device="cpu")

```

```

=====
Layer (type:depth-idx)                   Param #
=====
|---Conv2d: 1-1                           112
|---BatchNorm2d: 1-2                       8
|---ReLU: 1-3                             --
|---MaxPool2d: 1-4                         --
|---Conv2d: 1-5                           296
|---BatchNorm2d: 1-6                       16
|---ReLU: 1-7                             --
|---MaxPool2d: 1-8                         --
|---Flatten: 1-9                           --
|---Linear: 1-10                          131,136
|---ReLU: 1-11                             --
|---Linear: 1-12                          390
|---LogSoftmax: 1-13                       --
=====
Total params: 131,958
Trainable params: 131,958
Non-trainable params: 0
=====

```

```

Out[7]: =====
Layer (type:depth-idx)                Param #
=====
|Conv2d: 1-1                          112
|BatchNorm2d: 1-2                      8
|ReLU: 1-3                            --
|MaxPool2d: 1-4                       --
|Conv2d: 1-5                          296
|BatchNorm2d: 1-6                     16
|ReLU: 1-7                            --
|MaxPool2d: 1-8                       --
|Flatten: 1-9                         --
|Linear: 1-10                         131,136
|ReLU: 1-11                           --
|Linear: 1-12                         390
|LogSoftmax: 1-13                    --
=====
Total params: 131,958
Trainable params: 131,958
Non-trainable params: 0
=====

```

## Task 2

DataLoaders wrap around Datasets to provide efficient data batching, shuffling, and parallel loading during model training or inference. To define a custom dataset we must implement three functions: **init**, **len** and **get\_item**. While **len** defines the length of the dataset and thus the number of batches in the dataloader, **get\_item** can be used to get a single sample through the index.

```

In [8]: class Image_Dataset(Dataset):
        def __init__(self, X, Y):
            self.X = X
            self.Y = Y

        def __len__(self):
            # TODO: Return the length of the dataset
            return len(self.X)

        def __getitem__(self, idx):
            # TODO: Return the image as a float tensor and the label as a long t
            # Images need to be transposed from (H, W, C) to (C, H, W) for PyTor
            image = tc.tensor(self.X[idx].transpose(2, 0, 1), dtype=tc.float32)
            label = tc.tensor(self.Y[idx], dtype=tc.long)
            return image, label

```

```

In [9]: train_batch_size = 64
        val_batch_size = len(Y_val)
        test_batch_size = len(Y_test)

        # TODO: Create the dataset and dataloader
        train_dataset = Image_Dataset(X_train, Y_train)
        val_dataset = Image_Dataset(X_val, Y_val)
        test_dataset = Image_Dataset(X_test, Y_test)

        train_loader = DataLoader(train_dataset, batch_size=train_batch_size, shuffl


```

```
val_loader = DataLoader(val_dataset, batch_size=val_batch_size, shuffle=False)
test_loader = DataLoader(test_dataset, batch_size=test_batch_size, shuffle=False)
```

To make sure that everything has worked properly, we take a sample of the data\_loader and visualize it.

```
In [10]: sample_X, sample_Y = next(iter(train_loader))
plt.imshow(sample_X[0].T)
plt.title("y = " + str(int(sample_Y[0].item())))
plt.show()
```

```
/var/folders/gl/m711nqcx42d81f7zr656_v600000gn/T/ipykernel_1541/2461653407.
py:2: UserWarning: The use of `x.T` on tensors of dimension other than 2 to
reverse their shape is deprecated and it will throw an error in a future re
lease. Consider `x.mT` to transpose batches of matrices or `x.permute(*torc
h.arange(x.ndim - 1, -1, -1))` to reverse the dimensions of a tensor. (Trig
gered internally at /Users/runner/work/pytorch/pytorch/pytorch/aten/src/ATe
n/native/TensorShape.cpp:4416.)
  plt.imshow(sample_X[0].T)
```



### Task 3

Implement the training loop. Use the negative log-likelihood loss (NLLLoss) and the Adam optimizer. Be sure to zero the gradients after each optimization step to avoid accumulating contributions from previous epochs and batches.

```
In [13]: def train_cnn(model, train_loader, val_loader, lr, n_epochs, device):
    model = model.to(device)

    # TODO: Initialize the optimizer and loss function
    loss_function = nn.NLLLoss()
    optimizer = optim.Adam(model.parameters(), lr=lr)
```

```

train_loss = np.zeros(n_epochs)
val_loss = np.zeros(n_epochs)
train_acc = np.zeros(n_epochs)
val_acc = np.zeros(n_epochs)

for epoch in range(1, n_epochs + 1):
    model.train()

    epoch_loss = 0

    for X, Y in train_loader:
        X, Y = X.to(device), Y.to(device)
        # TODO: Implement the training loop
        optimizer.zero_grad()
        output = model(X)
        loss = loss_function(output, Y)
        loss.backward()
        optimizer.step()

    epoch_loss += loss.item()/len(train_loader)

    train_loss[epoch - 1] = epoch_loss
    train_acc[epoch - 1] = (output.argmax(dim=1) == Y).float().mean().item()

    model.eval()

    with tc.no_grad():
        X, Y = next(iter(val_loader))
        X, Y = X.to(device), Y.to(device)
        # TODO: Implement the evaluation step
        output = model(X)
        loss = loss_function(output, Y)

    val_loss[epoch - 1] = loss.item()
    val_acc[epoch - 1] = (output.argmax(dim=1) == Y).float().mean().item()
    print(f"Epoch {epoch}/{n_epochs} - Train Loss: {epoch_loss:.4f}, Test Loss: {val_loss:.4f}, Train Acc: {train_acc:.4f}, Val Acc: {val_acc:.4f}")

return train_loss, val_loss, train_acc, val_acc

```

```

In [19]: n_epochs = 50
# TODO: Train the model with different learning rates
learning_rates = [1e-1, 1e-3, 1e-4, 1e-5]
results = {}

for lr in learning_rates:
    print(f"\nTraining with learning rate: {lr}")
    model = CNN_Classifier()
    train_loss, val_loss, train_acc, val_acc = train_cnn(
        model, train_loader, val_loader, lr, n_epochs, device
    )
    results[lr] = {
        'model': model,
        'train_loss': train_loss,
        'val_loss': val_loss,
        'train_acc': train_acc,
        'val_acc': val_acc
    }

```

Training with learning rate: 0.1

Epoch 1/50 – Train Loss: 24.1252, Test Loss: 1.8230  
Epoch 2/50 – Train Loss: 1.8122, Test Loss: 2.5632  
Epoch 3/50 – Train Loss: 1.7964, Test Loss: 1.7932  
Epoch 4/50 – Train Loss: 1.7949, Test Loss: 1.7925  
Epoch 5/50 – Train Loss: 1.7954, Test Loss: 1.7921  
Epoch 6/50 – Train Loss: 1.7937, Test Loss: 1.7921  
Epoch 7/50 – Train Loss: 1.7929, Test Loss: 1.7920  
Epoch 8/50 – Train Loss: 1.7954, Test Loss: 1.7922  
Epoch 9/50 – Train Loss: 1.7939, Test Loss: 1.7926  
Epoch 10/50 – Train Loss: 1.7943, Test Loss: 1.7919  
Epoch 11/50 – Train Loss: 1.7937, Test Loss: 1.7921  
Epoch 12/50 – Train Loss: 1.7970, Test Loss: 1.7923  
Epoch 13/50 – Train Loss: 1.7968, Test Loss: 1.7943  
Epoch 14/50 – Train Loss: 1.7947, Test Loss: 1.7928  
Epoch 15/50 – Train Loss: 1.7987, Test Loss: 1.7925  
Epoch 16/50 – Train Loss: 1.7989, Test Loss: 1.7931  
Epoch 17/50 – Train Loss: 1.7977, Test Loss: 1.7933  
Epoch 18/50 – Train Loss: 1.7964, Test Loss: 1.7941  
Epoch 19/50 – Train Loss: 1.7954, Test Loss: 1.7923  
Epoch 20/50 – Train Loss: 1.7964, Test Loss: 1.7931  
Epoch 21/50 – Train Loss: 1.7951, Test Loss: 1.7925  
Epoch 22/50 – Train Loss: 1.7975, Test Loss: 1.7927  
Epoch 23/50 – Train Loss: 1.7972, Test Loss: 1.7931  
Epoch 24/50 – Train Loss: 1.7990, Test Loss: 1.7928  
Epoch 25/50 – Train Loss: 1.7964, Test Loss: 1.7927  
Epoch 26/50 – Train Loss: 1.7984, Test Loss: 1.7924  
Epoch 27/50 – Train Loss: 1.7955, Test Loss: 1.7934  
Epoch 28/50 – Train Loss: 1.7944, Test Loss: 1.7921  
Epoch 29/50 – Train Loss: 1.7954, Test Loss: 1.7922  
Epoch 30/50 – Train Loss: 1.7959, Test Loss: 1.7926  
Epoch 31/50 – Train Loss: 1.7952, Test Loss: 1.7923  
Epoch 32/50 – Train Loss: 1.7985, Test Loss: 1.7932  
Epoch 33/50 – Train Loss: 1.7981, Test Loss: 1.7926  
Epoch 34/50 – Train Loss: 1.7973, Test Loss: 1.7928  
Epoch 35/50 – Train Loss: 1.7980, Test Loss: 1.7929  
Epoch 36/50 – Train Loss: 1.7945, Test Loss: 1.7925  
Epoch 37/50 – Train Loss: 1.7952, Test Loss: 1.7922  
Epoch 38/50 – Train Loss: 1.7949, Test Loss: 1.7925  
Epoch 39/50 – Train Loss: 1.7963, Test Loss: 1.7920  
Epoch 40/50 – Train Loss: 1.7984, Test Loss: 1.7933  
Epoch 41/50 – Train Loss: 1.7975, Test Loss: 1.7931  
Epoch 42/50 – Train Loss: 1.7964, Test Loss: 1.7931  
Epoch 43/50 – Train Loss: 1.8002, Test Loss: 1.7926  
Epoch 44/50 – Train Loss: 1.7954, Test Loss: 1.7929  
Epoch 45/50 – Train Loss: 1.7953, Test Loss: 1.7922  
Epoch 46/50 – Train Loss: 1.7966, Test Loss: 1.7921  
Epoch 47/50 – Train Loss: 1.7955, Test Loss: 1.7936  
Epoch 48/50 – Train Loss: 1.7958, Test Loss: 1.7926  
Epoch 49/50 – Train Loss: 1.7955, Test Loss: 1.7919  
Epoch 50/50 – Train Loss: 1.7962, Test Loss: 1.7920

Training with learning rate: 0.001

Epoch 1/50 – Train Loss: 1.5027, Test Loss: 1.7479  
Epoch 2/50 – Train Loss: 0.8677, Test Loss: 2.1267  
Epoch 3/50 – Train Loss: 0.5151, Test Loss: 1.4809  
Epoch 4/50 – Train Loss: 0.3310, Test Loss: 0.8321  
Epoch 5/50 – Train Loss: 0.2247, Test Loss: 0.4975  
Epoch 6/50 – Train Loss: 0.1585, Test Loss: 0.6596  
Epoch 7/50 – Train Loss: 0.1109, Test Loss: 0.4652  
Epoch 8/50 – Train Loss: 0.0775, Test Loss: 0.3986



Epoch 9/50 – Train Loss: 0.0662, Test Loss: 0.3639  
Epoch 10/50 – Train Loss: 0.0493, Test Loss: 0.3199  
Epoch 11/50 – Train Loss: 0.0333, Test Loss: 0.3288  
Epoch 12/50 – Train Loss: 0.0271, Test Loss: 0.3127  
Epoch 13/50 – Train Loss: 0.0235, Test Loss: 0.3848  
Epoch 14/50 – Train Loss: 0.0193, Test Loss: 0.2970  
Epoch 15/50 – Train Loss: 0.0154, Test Loss: 0.2904  
Epoch 16/50 – Train Loss: 0.0134, Test Loss: 0.3685  
Epoch 17/50 – Train Loss: 0.0115, Test Loss: 0.3187  
Epoch 18/50 – Train Loss: 0.0098, Test Loss: 0.3201  
Epoch 19/50 – Train Loss: 0.0094, Test Loss: 0.3142  
Epoch 20/50 – Train Loss: 0.0081, Test Loss: 0.3089  
Epoch 21/50 – Train Loss: 0.0069, Test Loss: 0.3160  
Epoch 22/50 – Train Loss: 0.0065, Test Loss: 0.3368  
Epoch 23/50 – Train Loss: 0.0059, Test Loss: 0.3100  
Epoch 24/50 – Train Loss: 0.0051, Test Loss: 0.3254  
Epoch 25/50 – Train Loss: 0.0050, Test Loss: 0.3541  
Epoch 26/50 – Train Loss: 0.0044, Test Loss: 0.3155  
Epoch 27/50 – Train Loss: 0.0040, Test Loss: 0.3286  
Epoch 28/50 – Train Loss: 0.0038, Test Loss: 0.3138  
Epoch 29/50 – Train Loss: 0.0035, Test Loss: 0.3391  
Epoch 30/50 – Train Loss: 0.0032, Test Loss: 0.3292  
Epoch 31/50 – Train Loss: 0.0031, Test Loss: 0.3274  
Epoch 32/50 – Train Loss: 0.0027, Test Loss: 0.3349  
Epoch 33/50 – Train Loss: 0.0025, Test Loss: 0.3404  
Epoch 34/50 – Train Loss: 0.0024, Test Loss: 0.3488  
Epoch 35/50 – Train Loss: 0.0022, Test Loss: 0.3329  
Epoch 36/50 – Train Loss: 0.0021, Test Loss: 0.3436  
Epoch 37/50 – Train Loss: 0.0021, Test Loss: 0.3419  
Epoch 38/50 – Train Loss: 0.0020, Test Loss: 0.3432  
Epoch 39/50 – Train Loss: 0.0018, Test Loss: 0.3441  
Epoch 40/50 – Train Loss: 0.0018, Test Loss: 0.3329  
Epoch 41/50 – Train Loss: 0.0017, Test Loss: 0.3477  
Epoch 42/50 – Train Loss: 0.0016, Test Loss: 0.3502  
Epoch 43/50 – Train Loss: 0.0015, Test Loss: 0.3579  
Epoch 44/50 – Train Loss: 0.0014, Test Loss: 0.3431  
Epoch 45/50 – Train Loss: 0.0014, Test Loss: 0.3438  
Epoch 46/50 – Train Loss: 0.0013, Test Loss: 0.3587  
Epoch 47/50 – Train Loss: 0.0013, Test Loss: 0.3469  
Epoch 48/50 – Train Loss: 0.0012, Test Loss: 0.3517  
Epoch 49/50 – Train Loss: 0.0011, Test Loss: 0.3418  
Epoch 50/50 – Train Loss: 0.0011, Test Loss: 0.3639

Training with learning rate: 0.0001

Epoch 1/50 – Train Loss: 1.7334, Test Loss: 1.7869  
Epoch 2/50 – Train Loss: 1.5546, Test Loss: 1.7271  
Epoch 3/50 – Train Loss: 1.3683, Test Loss: 1.5448  
Epoch 4/50 – Train Loss: 1.1951, Test Loss: 1.3224  
Epoch 5/50 – Train Loss: 1.0445, Test Loss: 1.1279  
Epoch 6/50 – Train Loss: 0.9235, Test Loss: 1.0233  
Epoch 7/50 – Train Loss: 0.8341, Test Loss: 0.9288  
Epoch 8/50 – Train Loss: 0.7452, Test Loss: 0.8967  
Epoch 9/50 – Train Loss: 0.6831, Test Loss: 0.8141  
Epoch 10/50 – Train Loss: 0.6313, Test Loss: 0.7902  
Epoch 11/50 – Train Loss: 0.5789, Test Loss: 0.7342  
Epoch 12/50 – Train Loss: 0.5332, Test Loss: 0.7192  
Epoch 13/50 – Train Loss: 0.4916, Test Loss: 0.7040  
Epoch 14/50 – Train Loss: 0.4625, Test Loss: 0.6450  
Epoch 15/50 – Train Loss: 0.4298, Test Loss: 0.6246  
Epoch 16/50 – Train Loss: 0.4055, Test Loss: 0.6364  
Epoch 17/50 – Train Loss: 0.3709, Test Loss: 0.6077

Epoch 18/50 – Train Loss: 0.3521, Test Loss: 0.5807  
Epoch 19/50 – Train Loss: 0.3291, Test Loss: 0.5946  
Epoch 20/50 – Train Loss: 0.3072, Test Loss: 0.5136  
Epoch 21/50 – Train Loss: 0.2897, Test Loss: 0.5611  
Epoch 22/50 – Train Loss: 0.2726, Test Loss: 0.5471  
Epoch 23/50 – Train Loss: 0.2540, Test Loss: 0.4923  
Epoch 24/50 – Train Loss: 0.2396, Test Loss: 0.5357  
Epoch 25/50 – Train Loss: 0.2292, Test Loss: 0.4661  
Epoch 26/50 – Train Loss: 0.2144, Test Loss: 0.4804  
Epoch 27/50 – Train Loss: 0.2004, Test Loss: 0.4563  
Epoch 28/50 – Train Loss: 0.1872, Test Loss: 0.4384  
Epoch 29/50 – Train Loss: 0.1768, Test Loss: 0.4510  
Epoch 30/50 – Train Loss: 0.1665, Test Loss: 0.4464  
Epoch 31/50 – Train Loss: 0.1553, Test Loss: 0.4082  
Epoch 32/50 – Train Loss: 0.1487, Test Loss: 0.4207  
Epoch 33/50 – Train Loss: 0.1415, Test Loss: 0.4053  
Epoch 34/50 – Train Loss: 0.1312, Test Loss: 0.4066  
Epoch 35/50 – Train Loss: 0.1248, Test Loss: 0.4030  
Epoch 36/50 – Train Loss: 0.1164, Test Loss: 0.4082  
Epoch 37/50 – Train Loss: 0.1119, Test Loss: 0.3930  
Epoch 38/50 – Train Loss: 0.1070, Test Loss: 0.4247  
Epoch 39/50 – Train Loss: 0.0995, Test Loss: 0.3874  
Epoch 40/50 – Train Loss: 0.0951, Test Loss: 0.3513  
Epoch 41/50 – Train Loss: 0.0884, Test Loss: 0.3714  
Epoch 42/50 – Train Loss: 0.0843, Test Loss: 0.3510  
Epoch 43/50 – Train Loss: 0.0810, Test Loss: 0.3316  
Epoch 44/50 – Train Loss: 0.0774, Test Loss: 0.3399  
Epoch 45/50 – Train Loss: 0.0737, Test Loss: 0.3348  
Epoch 46/50 – Train Loss: 0.0697, Test Loss: 0.3343  
Epoch 47/50 – Train Loss: 0.0679, Test Loss: 0.3268  
Epoch 48/50 – Train Loss: 0.0633, Test Loss: 0.3227  
Epoch 49/50 – Train Loss: 0.0600, Test Loss: 0.3273  
Epoch 50/50 – Train Loss: 0.0573, Test Loss: 0.3080

Training with learning rate: 1e-05

Epoch 1/50 – Train Loss: 1.8014, Test Loss: 1.7932  
Epoch 2/50 – Train Loss: 1.7795, Test Loss: 1.7912  
Epoch 3/50 – Train Loss: 1.7627, Test Loss: 1.7782  
Epoch 4/50 – Train Loss: 1.7471, Test Loss: 1.7612  
Epoch 5/50 – Train Loss: 1.7318, Test Loss: 1.7459  
Epoch 6/50 – Train Loss: 1.7161, Test Loss: 1.7328  
Epoch 7/50 – Train Loss: 1.6995, Test Loss: 1.7202  
Epoch 8/50 – Train Loss: 1.6827, Test Loss: 1.7059  
Epoch 9/50 – Train Loss: 1.6662, Test Loss: 1.6912  
Epoch 10/50 – Train Loss: 1.6492, Test Loss: 1.6760  
Epoch 11/50 – Train Loss: 1.6314, Test Loss: 1.6594  
Epoch 12/50 – Train Loss: 1.6137, Test Loss: 1.6429  
Epoch 13/50 – Train Loss: 1.5953, Test Loss: 1.6251  
Epoch 14/50 – Train Loss: 1.5766, Test Loss: 1.6080  
Epoch 15/50 – Train Loss: 1.5594, Test Loss: 1.5910  
Epoch 16/50 – Train Loss: 1.5405, Test Loss: 1.5736  
Epoch 17/50 – Train Loss: 1.5220, Test Loss: 1.5576  
Epoch 18/50 – Train Loss: 1.5039, Test Loss: 1.5398  
Epoch 19/50 – Train Loss: 1.4848, Test Loss: 1.5226  
Epoch 20/50 – Train Loss: 1.4661, Test Loss: 1.5052  
Epoch 21/50 – Train Loss: 1.4479, Test Loss: 1.4891  
Epoch 22/50 – Train Loss: 1.4287, Test Loss: 1.4725  
Epoch 23/50 – Train Loss: 1.4101, Test Loss: 1.4565  
Epoch 24/50 – Train Loss: 1.3911, Test Loss: 1.4397  
Epoch 25/50 – Train Loss: 1.3718, Test Loss: 1.4228  
Epoch 26/50 – Train Loss: 1.3525, Test Loss: 1.4056

```

Epoch 27/50 - Train Loss: 1.3337, Test Loss: 1.3880
Epoch 28/50 - Train Loss: 1.3151, Test Loss: 1.3701
Epoch 29/50 - Train Loss: 1.2963, Test Loss: 1.3535
Epoch 30/50 - Train Loss: 1.2793, Test Loss: 1.3377
Epoch 31/50 - Train Loss: 1.2607, Test Loss: 1.3212
Epoch 32/50 - Train Loss: 1.2428, Test Loss: 1.3046
Epoch 33/50 - Train Loss: 1.2248, Test Loss: 1.2880
Epoch 34/50 - Train Loss: 1.2073, Test Loss: 1.2726
Epoch 35/50 - Train Loss: 1.1909, Test Loss: 1.2575
Epoch 36/50 - Train Loss: 1.1737, Test Loss: 1.2426
Epoch 37/50 - Train Loss: 1.1570, Test Loss: 1.2283
Epoch 38/50 - Train Loss: 1.1415, Test Loss: 1.2152
Epoch 39/50 - Train Loss: 1.1255, Test Loss: 1.2013
Epoch 40/50 - Train Loss: 1.1100, Test Loss: 1.1874
Epoch 41/50 - Train Loss: 1.0952, Test Loss: 1.1728
Epoch 42/50 - Train Loss: 1.0806, Test Loss: 1.1631
Epoch 43/50 - Train Loss: 1.0661, Test Loss: 1.1481
Epoch 44/50 - Train Loss: 1.0515, Test Loss: 1.1359
Epoch 45/50 - Train Loss: 1.0369, Test Loss: 1.1232
Epoch 46/50 - Train Loss: 1.0242, Test Loss: 1.1123
Epoch 47/50 - Train Loss: 1.0107, Test Loss: 1.0999
Epoch 48/50 - Train Loss: 0.9981, Test Loss: 1.0892
Epoch 49/50 - Train Loss: 0.9854, Test Loss: 1.0797
Epoch 50/50 - Train Loss: 0.9729, Test Loss: 1.0696

```

```

In [20]: # TODO: Visualize the results
fig, axes = plt.subplots(2, 2, figsize=(15, 12))

for lr in learning_rates:
    epochs = range(1, n_epochs + 1)

    # Plot losses
    axes[0, 0].plot(epochs, results[lr]['train_loss'], label=f'Train LR={lr}')
    axes[0, 1].plot(epochs, results[lr]['val_loss'], label=f'Val LR={lr}')

    # Plot accuracies
    axes[1, 0].plot(epochs, results[lr]['train_acc'], label=f'Train LR={lr}')
    axes[1, 1].plot(epochs, results[lr]['val_acc'], label=f'Val LR={lr}')

axes[0, 0].set_title('Training Loss')
axes[0, 0].set_xlabel('Epoch')
axes[0, 0].set_ylabel('Loss')
axes[0, 0].legend()
axes[0, 0].grid(True)

axes[0, 1].set_title('Validation Loss')
axes[0, 1].set_xlabel('Epoch')
axes[0, 1].set_ylabel('Loss')
axes[0, 1].legend()
axes[0, 1].grid(True)


axes[1, 0].set_title('Training Accuracy')
axes[1, 0].set_xlabel('Epoch')
axes[1, 0].set_ylabel('Accuracy')
axes[1, 0].legend()
axes[1, 0].grid(True)

axes[1, 1].set_title('Validation Accuracy')
axes[1, 1].set_xlabel('Epoch')
axes[1, 1].set_ylabel('Accuracy')
axes[1, 1].legend()

```

```
axes[1, 1].grid(True)
```

```
plt.tight_layout()
plt.show()
```



- A learning rate of 0.001 seems to be the most effective among the tested values, providing a good balance between convergence speed and generalization performance.
- A learning rate of 0.1 is too high, leading to poor performance.
- Learning rates of 0.0001 and 0.00001 are too low, resulting in slow convergence.
- The model is able to achieve a validation accuracy of around 90% with the optimal learning rate of 0.001, indicating good generalization to unseen data.

## Task 4

```
In [21]: # TODO: apply the best model to the test set
best_lr = max(learning_rates, key=lambda lr: results[lr]['val_acc'][-1])
best_model = results[best_lr]['model']

print(f"\nBest learning rate: {best_lr}")
print(f"Best validation accuracy: {results[best_lr]['val_acc'][-1]:.4f}")

# Test the best model
best_model.eval()
with tc.no_grad():
    X_test_tensor, Y_test_tensor = next(iter(test_loader))
    X_test_tensor, Y_test_tensor = X_test_tensor.to(device), Y_test_tensor.to(device)
```

```
test_output = best_model(X_test_tensor)
test_acc = (test_output.argmax(dim=1) == Y_test_tensor).float().mean().item()

print(f"Test accuracy: {test_acc:.4f}")
print(f"Difference from validation accuracy: {abs(test_acc - results[best_val_acc]):.4f}")
```

Best learning rate: 0.001

Best validation accuracy: 0.9083

Test accuracy: 0.8917

Difference from validation accuracy: 0.0167